

```
In [1]: # Generic inputs for most ML tasks
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
from sklearn.ensemble import RandomForestRegressor

pd.options.display.float_format = '{:,.2f}'.format

# setup interactive notebook mode
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"

from IPython.display import display, HTML
```

Read and pre-process data

```
In [2]: # fetch data

diamonds_data = pd.read_csv('diamonds.csv')

diamonds_data.head()
```

```
Out [2]:
```

	rownames	carat	cut	color	clarity	depth	table	price	x	y	z
0	1	0.23	Ideal	E	SI2	61.50	55.00	326	3.95	3.98	2.43
1	2	0.21	Premium	E	SI1	59.80	61.00	326	3.89	3.84	2.31
2	3	0.23	Good	E	VS1	56.90	65.00	327	4.05	4.07	2.31
3	4	0.29	Premium	I	VS2	62.40	58.00	334	4.20	4.23	2.63
4	5	0.31	Good	J	SI2	63.30	58.00	335	4.34	4.35	2.75

```
In [3]: #column data types
diamonds_data.dtypes
```

```
Out [3]: rownames      int64
carat      float64
cut        object
color      object
clarity     object
depth      float64
table      float64
price      int64
x          float64
y          float64
z          float64
dtype: object
```

In [4]: *#dropping rownames column*

```
diamonds_data = diamonds_data.drop(columns = ['rownames'], axis=1)
diamonds_data.head()
```

Out [4]:

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.50	55.00	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.80	61.00	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.90	65.00	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.40	58.00	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.30	58.00	335	4.34	4.35	2.75

In [5]: *#NaN values*

```
diamonds_data.isna().sum()
```

Out [5]:

carat	0
cut	0
color	0
clarity	0
depth	0
table	0
price	0
x	0
y	0
z	0
dtype:	int64

There are no NaN values in the dataframe

In [80]: *#No. of rows*

```
print("Number of rows in the dataframe are",diamonds_data.shape[0])
```

Number of rows in the dataframe are 53940

In [7]: *#values taken by the three categorical variables*

```
cut_values = diamonds_data['cut'].unique()
print("Set of values for 'cut':", cut_values)

# To get unique values for the 'color' categorical variable
color_values = diamonds_data['color'].unique()
print("Set of values for 'color':", color_values)

# To get unique values for the 'clarity' categorical variable
clarity_values = diamonds_data['clarity'].unique()
print("Set of values for 'clarity':", clarity_values)
```

```
Set of values for 'cut': ['Ideal' 'Premium' 'Good' 'Very Good' 'Fair']
Set of values for 'color': ['E' 'I' 'J' 'H' 'F' 'G' 'D']
Set of values for 'clarity': ['SI2' 'SI1' 'VS1' 'VS2' 'VVS2' 'VVS1' 'I1' 'IF']
```

Simple Linear Regression

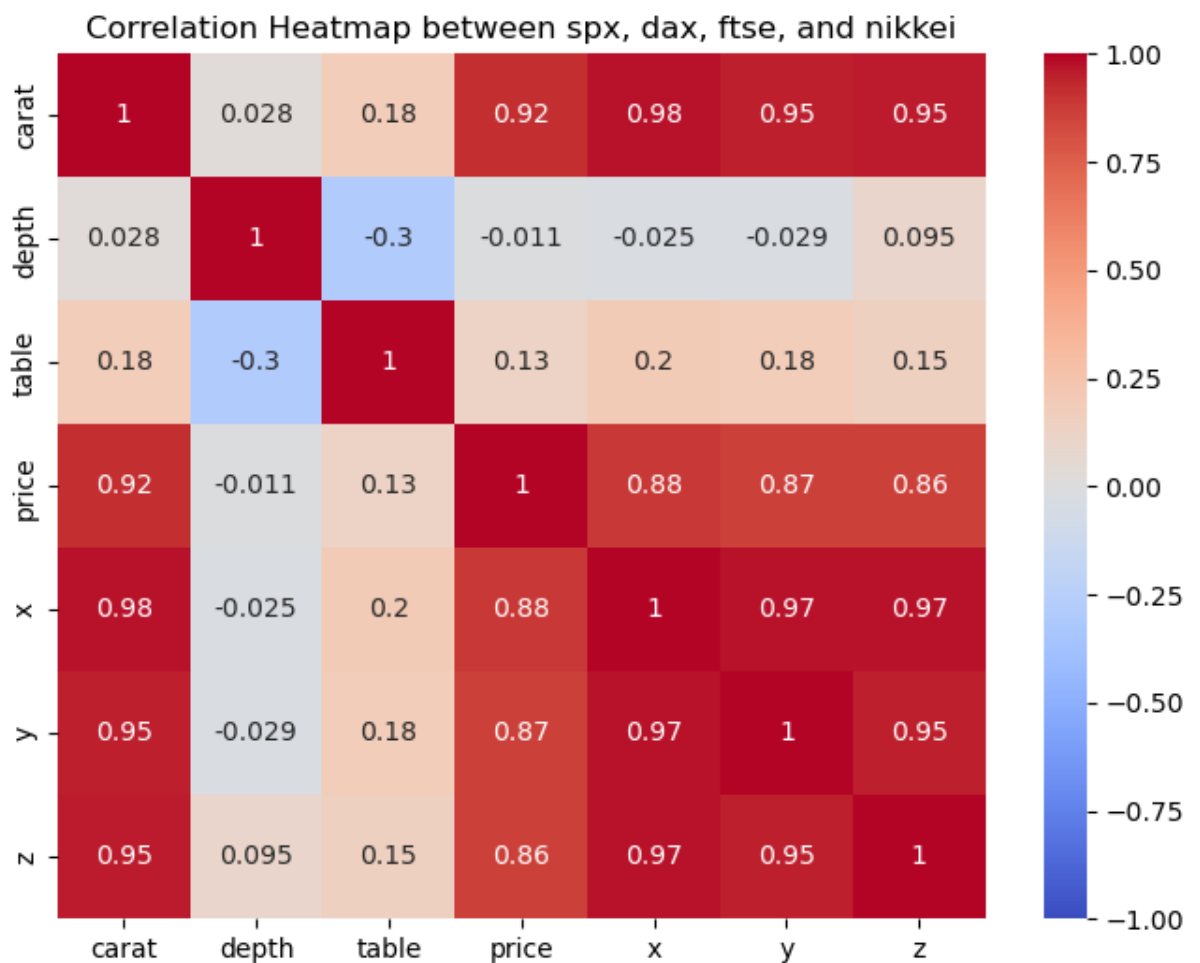
```
In [8]: import seaborn as sns
df_without_categorical_values = diamonds_data.drop(columns = ['cut', 'color', 'clarity'])
corr_matrix = df_without_categorical_values.corr()

# Create a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Heatmap between spx, dax, ftse, and nikkei')
plt.show()
```

Out[8]: <Figure size 800x600 with 0 Axes>

Out[8]: <Axes: >

Out[8]: Text(0.5, 1.0, 'Correlation Heatmap between spx, dax, ftse, and nikkei')



carat has the highest correlation with price

```
In [9]: X = diamonds_data['carat']
```

```
In [10]: Y = diamonds_data['price']
```

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.25)
```

```
In [12]: train_length = X_train.shape[0]
test_length = X_test.shape[0]

print('Length of train and test data are:', train_length , test_length )

Length of train and test data are: 40455 13485
```

```
In [13]: first_row_index_train = X_train.index[0]
first_row_index_test = X_test.index[0]

print("First row index of X_train:", first_row_index_train)
print("First row index of X_test:", first_row_index_test)

First row index of X_train: 32301
First row index of X_test: 44751
```

```
In [14]: model = LinearRegression(fit_intercept = True)
model.fit(X_train.array.reshape(-1, 1), y_train.array.reshape(-1, 1))
```

```
Out[14]: ▾ LinearRegression
LinearRegression()
```

```
In [15]: # The following gives the R-square score
model.score(X_train.array.reshape(-1, 1), y_train.array.reshape(-1, 1))

# This is the coefficient Beta_1 (or slope of the Simple Linear Regression)
model.coef_

# This is the coefficient Beta_0
model.intercept_
```

```
Out[15]: 0.8486051211546091
```

```
Out[15]: array([[7740.34546342]])
```

```
Out[15]: array([-2244.80400041])
```

An R-square score of approximately 0.85 suggests that the model explains about 85% of the variance in the target variable, which is generally considered a strong correlation and indicates a good fit.

Coefficients

1. Beta_1 (Coefficient for the Independent Variable): 7740.34546342 This coefficient represents the slope of the regression line in the case of simple linear regression (or the effect size of the variable in multiple regression). It indicates how much the dependent variable is expected to increase (or decrease, if the coefficient was negative) when the independent variable increases by one unit. In this context, a coefficient of approximately

7740.35 means that for each one-unit increase in the independent variable, the dependent variable is expected to increase by about 7740.35 units. This value suggests a strong positive relationship between the independent variable and the dependent variable.

2. Beta_0 (Intercept): -2244.80400041 The intercept, or Beta_0, represents the expected value of the dependent variable when all independent variables are equal to zero. An intercept of

```
In [250]: # Calculate R-squared on test data
model.score(X_test.array.reshape(-1, 1), y_test.array.reshape(-1, 1))
```

```
Out[250]: 0.8514897700970667
```

```
In [84]: test_output = pd.DataFrame(model.predict(X_test.array.reshape(-1, 1)), id
```

```
In [17]: test_output.head()
```

```
Out[17]:
```

	pred_price
44751	1,702.77
22963	13,313.29
9078	6,114.77
26148	14,242.13
29451	309.51

```
In [18]: test_output = test_output.merge(y_test, left_index = True, right_index =
test_output.head()
mean_absolute_error = abs(test_output['pred_price'] - test_output['price
print('Mean absolute error is ')
print(mean_absolute_error)
```

```
Out[18]:
```

	pred_price	price
44751	1,702.77	1619
22963	13,313.29	11011
9078	6,114.77	4521
26148	14,242.13	15454
29451	309.51	702

Mean absolute error is
998.4973200867355

```
In [19]: average_price_test = y_test.mean()
```

```
In [20]: fraction_mae = mean_absolute_error / average_price_test
print("Fraction of MAE to Average Price:", fraction_mae)
```

Fraction of MAE to Average Price: 0.2541206399951147

A fraction of approximately 0.254 means that the mean absolute error is about 25.4% of the average price in the test set.

An MAE of 998.5 and its fraction of 25.4% of the average price indicate that, despite the good fit, the average prediction error may still represent a significant portion of the price.

Multiple Linear Regression:

```
In [33]: X2 = diamonds_data.drop(columns = ['cut', 'color', 'clarity', 'price'],
```

```
In [34]: X2.head()
```

Out[34]:

	carat	depth	table	x	y	z
0	0.23	61.50	55.00	3.95	3.98	2.43
1	0.21	59.80	61.00	3.89	3.84	2.31
2	0.23	56.90	65.00	4.05	4.07	2.31
3	0.29	62.40	58.00	4.20	4.23	2.63
4	0.31	63.30	58.00	4.34	4.35	2.75

```
In [35]: Y2 = diamonds_data["price"]
```

```
In [36]: X2_train, X2_test, y2_train, y2_test = train_test_split(X2, Y2, test_size=
```

```
In [42]: X2_train
X2_test
y2_test
y2_train
```

Out [42]:

	carat	depth	table	x	y	z
32301	0.37	60.70	60.00	4.65	4.68	2.83
39009	0.40	61.70	57.00	4.77	4.73	2.93
22757	1.02	61.40	58.00	6.46	6.43	3.96
15129	1.07	62.30	55.00	6.59	6.54	4.09
17861	1.19	61.70	56.00	6.78	6.81	4.19
...
48417	0.70	62.40	58.00	5.62	5.67	3.52
22637	2.00	60.30	56.00	8.27	8.16	4.94
42891	0.51	62.80	57.00	5.12	5.10	3.21
38368	0.53	63.80	57.00	5.10	5.12	3.26
14000	1.19	62.70	61.00	6.73	6.66	4.20

40455 rows × 6 columns

Out [42]:

	carat	depth	table	x	y	z
44751	0.51	61.40	58.00	5.13	5.09	3.14
22963	2.01	62.90	54.00	8.06	7.93	5.05
9078	1.08	62.10	59.00	6.57	6.53	4.07
26148	2.13	61.50	57.00	8.27	8.34	5.11
29451	0.33	61.90	56.00	4.46	4.49	2.77
...
8884	1.04	62.40	56.00	6.40	6.48	4.02
34947	0.30	61.90	60.00	4.27	4.29	2.65
21821	1.50	60.20	61.00	7.27	7.32	4.39
9942	1.01	61.10	56.00	6.44	6.48	3.95
8296	1.13	61.70	57.00	6.70	6.59	4.13

13485 rows × 6 columns


```
Out[42]: 44751      1619
          22963      11011
          9078       4521
          26148     15454
          29451       702
          ...
          8884       4486
          34947       378
          21821      9892
          9942       4693
          8296       4385
          Name: price, Length: 13485, dtype: int64
```

```
Out[42]: 32301       454
          39009      1056
          22757     10773
          15129      6082
          17861      7207
          ...
          48417      1971
          22637     10685
          42891      1359
          38368      1023
          14000      5698
          Name: price, Length: 40455, dtype: int64
```

```
In [43]: train_length2 = X2_train.shape[0]
          test_length2 = X2_test.shape[0]

          print('Length of train and test data are:', train_length2 , test_length2)

          Length of train and test data are: 40455 13485
```

The length of test and train data are same as the lengths of SLR

```
In [44]: first_row_index_train2 = X2_train.index[0]
          first_row_index_test2 = X2_test.index[0]

          print("First row index of X_train:", first_row_index_train2)
          print("First row index of X_test:", first_row_index_test2)

          First row index of X_train: 32301
          First row index of X_test: 44751
```

First row index of both test and train data are same as that of SLR

```
In [45]: model2 = LinearRegression(fit_intercept = True)
          model2.fit(X2_train, y2_train)
```

```
Out[45]: ▾ LinearRegression
          LinearRegression()
```

```
In [46]: # The following gives the R-square score on train data
model2.score(X2_train, y2_train)

# This is the coefficient Beta_1, ..., Beta_7
model2.coef_

# This is the coefficient Beta_0
model2.intercept_
```

Out[46]: 0.8580892707297321

Out[46]: array([10572.42070164, -212.30889735, -100.98829263, -1339.48535974,
28.72790744, 207.71501298])

Out[46]: 21184.40695343088

The R-square score for model2 is approximately 0.858, which means that about 85.8% of the variance in the dependent variable is explained by the model. This is a strong score that suggests a good fit of the model to the data, especially in the context of multiple regression where more than one independent variable is used to predict the outcome.

```
In [251]: # The following gives the R-square score on train data
model2.score(X2_test, y2_test)
```

Out[251]: 0.8623985511113466

```
In [57]: test_output2 = pd.DataFrame(model2.predict(X2_test), index = X2_test.ind
```

```
In [58]: test_output2.head()
```

Out[58]:

	pred_price
44751	1,610.14
22963	14,107.90
9078	5,692.50
26148	15,113.80
29451	606.29

```
In [59]: test_output2 = test_output2.merge(y2_test, left_index = True, right_index=True)
test_output2.head()
mean_absolute_error2 = abs(test_output2['pred_price'] - test_output2['price'])
print('Mean absolute error is ')
print(mean_absolute_error2)
```

Out [59]:

	pred_price	price
44751	1,610.14	1619
22963	14,107.90	11011
9078	5,692.50	4521
26148	15,113.80	15454
29451	606.29	702

Mean absolute error is
882.599634844387

```
In [60]: average_price_test2 = y2_test.mean()
```

```
In [61]: fraction_mae2 = mean_absolute_error2 / average_price_test2
print("Fraction of MAE to Average Price:", fraction_mae2)
```

Fraction of MAE to Average Price: 0.224624322523597

This fraction indicates that the mean absolute error is about 22.5% of the average price in the test set.

Comparison with SLR For the SLR model, the reported MAE was approximately 998.5, and the fraction of MAE to the average price was about 25.4%. Compared to these values, the MLR model has shown an improvement in both the absolute size of the error (a lower MAE) and the error relative to the average price (a lower fraction).

Conclusion Yes, the predictions in the test set have improved with the MLR model compared to the SLR model. The reduction in both the MAE and its fraction to the average price indicates that incorporating multiple features into the model has enhanced its predictive accuracy, leading to closer predictions to the actual values and a reduction in the average error magnitude relative to the price scale.

Multiple Linear Regression with Categorical Values

```
In [197]: X3 = diamonds_data.drop(columns = ['price'], axis=1)
X3.head()
```

Out[197]:

	carat	cut	color	clarity	depth	table	x	y	z
0	0.23	Ideal	E	SI2	61.50	55.00	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.80	61.00	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.90	65.00	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.40	58.00	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.30	58.00	4.34	4.35	2.75

```
In [198]: Y3 = diamonds_data['price']
Y3.head()
```

Out[198]:

0	326
1	326
2	327
3	334
4	335

Name: price, dtype: int64

```
In [199]: print(diamonds_data['cut'].unique())
print(diamonds_data['color'].unique())
print(diamonds_data['clarity'].unique())

['Ideal' 'Premium' 'Good' 'Very Good' 'Fair']
['E' 'I' 'J' 'H' 'F' 'G' 'D']
['SI2' 'SI1' 'VS1' 'VS2' 'VVS2' 'VVS1' 'I1' 'IF']
```

```
In [200]: diamonds_data['cut'].value_counts()
diamonds_data['color'].value_counts()
diamonds_data['clarity'].value_counts()
```

```
Out[200]: cut
Ideal      21551
Premium    13791
Very Good  12082
Good       4906
Fair       1610
Name: count, dtype: int64
```

```
Out[200]: color
G      11292
E      9797
F      9542
H      8304
D      6775
I      5422
J      2808
Name: count, dtype: int64
```

```
Out[200]: clarity
SI1      13065
VS2      12258
SI2       9194
VS1       8171
VVS2       5066
VVS1       3655
IF         1790
I1         741
Name: count, dtype: int64
```

```
In [201]: from sklearn.preprocessing import OneHotEncoder

def get_ohc(df, col):
    ohe = OneHotEncoder(drop='first', handle_unknown='error', sparse_output=False)
    ohe.fit(df[[col]])
    temp_df = pd.DataFrame(data=ohe.transform(df[[col]]).toarray(), columns=ohe.get_feature_names_out([col])
    # If you have a newer version, replace with columns=ohe.get_feature_names_out([col])
    df.drop(columns=[col], axis=1, inplace=True)
    df = pd.concat([df.reset_index(drop=True), temp_df], axis=1)
    return df
```

```
In [202]: X3 = get_ohe(X3, 'cut')
X3 = get_ohe(X3, 'color')
X3 = get_ohe(X3, 'clarity')

X3.head(20)
```

Out[202]:

	carat	depth	table	x	y	z	cut_Good	cut_Ideal	cut_Premium	cut_Very Good	...	color_H
0	0.23	61.50	55.00	3.95	3.98	2.43	0	1	0	0	...	C
1	0.21	59.80	61.00	3.89	3.84	2.31	0	0	1	0	...	C
2	0.23	56.90	65.00	4.05	4.07	2.31	1	0	0	0	...	C
3	0.29	62.40	58.00	4.20	4.23	2.63	0	0	1	0	...	C
4	0.31	63.30	58.00	4.34	4.35	2.75	1	0	0	0	...	C
5	0.24	62.80	57.00	3.94	3.96	2.48	0	0	0	1	...	C
6	0.24	62.30	57.00	3.95	3.98	2.47	0	0	0	1	...	C
7	0.26	61.90	55.00	4.07	4.11	2.53	0	0	0	1	...	1
8	0.22	65.10	61.00	3.87	3.78	2.49	0	0	0	0	...	C
9	0.23	59.40	61.00	4.00	4.05	2.39	0	0	0	1	...	1
10	0.30	64.00	55.00	4.25	4.28	2.73	1	0	0	0	...	C
11	0.23	62.80	56.00	3.93	3.90	2.46	0	1	0	0	...	C
12	0.22	60.40	61.00	3.88	3.84	2.33	0	0	1	0	...	C
13	0.31	62.20	54.00	4.35	4.37	2.71	0	1	0	0	...	C
14	0.20	60.20	62.00	3.79	3.75	2.27	0	0	1	0	...	C
15	0.32	60.90	58.00	4.38	4.42	2.68	0	0	1	0	...	C
16	0.30	62.00	54.00	4.31	4.34	2.68	0	1	0	0	...	C
17	0.30	63.40	54.00	4.23	4.29	2.70	1	0	0	0	...	C
18	0.30	63.80	56.00	4.23	4.26	2.71	1	0	0	0	...	C
19	0.30	62.70	59.00	4.21	4.27	2.66	0	0	0	1	...	C

20 rows × 23 columns

Here's how we determine the number of columns in the eventual set after one-hot encoding:

let's consider the unique categories in each of the encoded categorical columns minus one (since we're dropping the first column):

'cut': If it originally has 5 unique values (Fair, Good, Very Good, Premium, Ideal), one-hot encoding with drop='first' will result in 4 columns. 'color': If it has 7 unique values (D through J), one-hot encoding with drop='first' will result in 6 columns. 'clarity': If it has 8 unique values (I1, SI2, SI1, VS2, VS1, VVS2, VVS1, IF), one-hot encoding with drop='first' will result in 7 columns.

Total Columns After Encoding Original non-encoded columns: 6 Encoded columns for 'cut': 4
Encoded columns for 'color': 6 Encoded columns for 'clarity': 7 Total = 6 (original) + 4 (cut) + 6
(color) + 7 (clarity) = 23 columns

... ..

```
In [204]: X3_train, X3_test, y3_train, y3_test = train_test_split(X3, Y3, test_size=
```

```
In [205]: model3 = LinearRegression(fit_intercept = True)
model3.fit(X3_train, y3_train)
```

```
Out[205]: ▼ LinearRegression
LinearRegression()
```

```
In [206]: # The following gives the R-square score on train data
model3.score(X3_train, y3_train)

# This is the coefficient Beta_1, ..., Beta_26
model3.coef_

# This is the coefficient Beta_0
model3.intercept_
```

```
Out[206]: 0.9197535054831985
```

```
Out[206]: array([ 1.11814801e+04, -5.85181599e+01, -2.43103736e+01, -9.47154791e+
02,
               4.71160598e+00, -9.21343078e+01,  5.64156045e+02,  8.35219660e+
02,
               7.65684380e+02,  7.33319794e+02, -2.25861063e+02, -2.76687793e+
02,
              -4.79782911e+02, -9.73649105e+02, -1.46411040e+03, -2.39532202e+
03,
               5.45238205e+03,  3.78052359e+03,  2.82037545e+03,  4.70977853e+
03,
               4.38909093e+03,  5.12101322e+03,  5.07443217e+03])
```

```
Out[206]: 1499.3392254312134
```

The enhanced multiple linear regression model, incorporating extra features like categorical data, significantly improves our result interpretation compared to its predecessor. It achieves a score close to 0.920, signifying a robust fit. The intercept, Beta_0, is estimated to be approximately 1499.339. Each of the model's coefficients indicates the degree to which a particular feature affects the outcome variable. This advancement provides a deeper insight into the various elements affecting predictions, thereby increasing the model's overall usefulness.

```
In [252]: # The following gives the R-square score on test data
model3.score(X3_test, y3_test)
```

```
Out[252]: 0.9197839099301064
```

```
In [207]: test_output3 = pd.DataFrame(model3.predict(X3_test), index = X3_test.index)
```

```
In [208]: test_output3.head()
```

Out[208]:

	pred_price
44751	1,952.74
22963	12,178.85
9078	5,091.65
26148	15,140.76
29451	137.09

```
In [209]: test_output3 = test_output3.merge(y3_test, left_index = True, right_index = True)
test_output3.head()
mean_absolute_error3 = abs(test_output3['pred_price'] - test_output3['price']).mean()
print('Mean absolute error is ')
print(mean_absolute_error3)
```

Out[209]:

	pred_price	price
44751	1,952.74	1619
22963	12,178.85	11011
9078	5,091.65	4521
26148	15,140.76	15454
29451	137.09	702

Mean absolute error is
744.8276400124315

An MAE of approximately 744.82 indicates that, on average, the model's predictions are about 744.82 units away from the actual values.

```
In [210]: average_price_test3 = y3_test.mean()
```

```
In [211]: fraction_mae3 = mean_absolute_error3 / average_price_test3
print("Fraction of MAE to Average Price:", fraction_mae3)
```

Fraction of MAE to Average Price: 0.18956092596179258

An average absolute error of approximately 745, coupled with its ratio to the mean price being roughly 0.19, indicates that our enhanced model, which includes additional variables, performs commendably in forecasting results. Given that the model now aligns more closely with the data, evidenced by an R-square value close to 0.92, it logically follows that there has been a reduction in predictive errors. This marks a significant enhancement from the previous error ratio of around 0.23, highlighting the beneficial impact of incorporating more variables. Essentially, it signifies an improvement in the predictive accuracy of our model.

Quantile Regression with Categorical Variables

```
In [212]: X4 = diamonds_data.drop(columns = ['price'], axis=1)
X4.head()
```

Out[212]:

	carat	cut	color	clarity	depth	table	x	y	z
0	0.23	Ideal	E	SI2	61.50	55.00	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.80	61.00	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.90	65.00	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.40	58.00	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.30	58.00	4.34	4.35	2.75

```
In [213]: Y4 = diamonds_data['price']
Y4.head()
```

Out[213]:

0	326
1	326
2	327
3	334
4	335

Name: price, dtype: int64

```
In [214]: X4 = get_ohe(X4, 'cut')
X4 = get_ohe(X4, 'color')
X4 = get_ohe(X4, 'clarity')

X4.head()
```

Out[214]:

	carat	depth	table	x	y	z	cut_Good	cut_Ideal	cut_Premium	cut_Very Good	...	color_H
0	0.23	61.50	55.00	3.95	3.98	2.43	0	1	0	0	...	0
1	0.21	59.80	61.00	3.89	3.84	2.31	0	0	1	0	...	0
2	0.23	56.90	65.00	4.05	4.07	2.31	1	0	0	0	...	0
3	0.29	62.40	58.00	4.20	4.23	2.63	0	0	1	0	...	0
4	0.31	63.30	58.00	4.34	4.35	2.75	1	0	0	0	...	0

5 rows × 23 columns

```
In [215]: X4_train, X4_test, y4_train, y4_test = train_test_split(X4, Y4, test_size=0.2)
```

```
In [217]: import statsmodels.formula.api as smf
import statsmodels.api as sm
```

```
In [253]: # Create the quantile regression model
mod = sm.QuantReg(Y4, X4)
# Fit the model using the desired quantile
res = mod.fit(q=0.5) # For median (50th percentile)

# Display the results
print(res.summary())
```

QuantReg Regression Results

```

=====
=====
Dep. Variable:          price    Pseudo R-squared:
0.7709
Model:                 QuantReg  Bandwidth:
54.45
Method:                Least Squares  Sparsity:
946.8
Date:                  Mon, 12 Feb 2024  No. Observations:
53940
Time:                  22:41:29    Df Residuals:
53917
                                   Df Model:
23
=====
=====

```

	coef	std err	t	P> t	[0.025
0.975]					
carat	1.293e+04	19.211	672.807	0.000	1.29e+04
1.3e+04					
depth	55.0542	0.839	65.585	0.000	53.409
56.699					
table	-20.2867	0.898	-22.591	0.000	-22.047
-18.527					
x	-559.4452	12.904	-43.353	0.000	-584.738
-534.152					
y	131.1420	8.078	16.234	0.000	115.308
146.976					
z	-2586.5157	13.423	-192.698	0.000	-2612.824
-2560.207					
cut_Good	358.7946	13.718	26.155	0.000	331.907
385.682					
cut_Ideal	520.4984	12.547	41.485	0.000	495.907
545.090					
cut_Premium	504.3973	12.709	39.688	0.000	479.487
529.307					
cut_Very Good	423.3799	12.733	33.251	0.000	398.424
448.336					
color_E	-138.8535	7.493	-18.530	0.000	-153.541
-124.167					
color_F	-205.8204	7.578	-27.160	0.000	-220.673
-190.967					
color_G	-271.0289	7.421	-36.520	0.000	-285.575
-256.483					
color_H	-482.3621	7.891	-61.132	0.000	-497.828
-466.897					
color_I	-842.1548	8.865	-94.997	0.000	-859.530
-824.779					
color_J	-1632.7617	10.946	-149.160	0.000	-1654.217
-1611.307					
clarity_IF	3311.2173	21.181	156.328	0.000	3269.702
3352.733					
clarity_SI1	2582.9559	18.183	142.054	0.000	2547.317
2618.594					

clarity_SI2	1899.7364	18.250	104.095	0.000	1863.966
1935.507					
clarity_VS1	3040.0118	18.526	164.098	0.000	3003.702
3076.322					
clarity_VS2	2910.6472	18.250	159.484	0.000	2874.876
2946.418					
clarity_VVS1	3197.3169	19.584	163.260	0.000	3158.932
3235.702					
clarity_VVS2	3168.3902	19.055	166.274	0.000	3131.042
3205.739					

=====

=====

The condition number is large, 1.99e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [259]: # Calculate predicted prices
# Assuming 'res' is the result object from fitting your quantile regress.
predicted_prices = X4_test.dot(res.params)
```

```
In [262]: import pandas as pd

# Create a DataFrame with predictions
test_output4 = pd.DataFrame(predicted_prices, columns=['pred_price'])
test_output4 = test_output4.merge(y4_test, left_index = True, right_index = True)
test_output4.head()
```

Out[262]:

	pred_price	price
44751	1,680.79	1619
22963	12,603.92	11011
9078	5,096.39	4521
26148	15,062.74	15454
29451	430.84	702

```
In [267]: mean_absolute_error4 = abs(test_output4['pred_price'] - test_output4['price'])
print('Mean absolute error is ')
print(mean_absolute_error4)
```

Mean absolute error is
644.5251556319222

```
In [268]: average_price_test4 = y4_test.mean()
```

```
In [269]: fraction_mae4 = mean_absolute_error4 / average_price_test4
print("Fraction of MAE to Average Price:", fraction_mae4)
```

Fraction of MAE to Average Price: 0.1640336350906855

